


# CrossEntropy Loss

How to train a classifier

# Classifier


A **classifier** is a type of **model** that is **trained** to **categorize** input data into **classes**.

The **most used loss** function to **train** a **classifier** is the **crossentropy** loss.



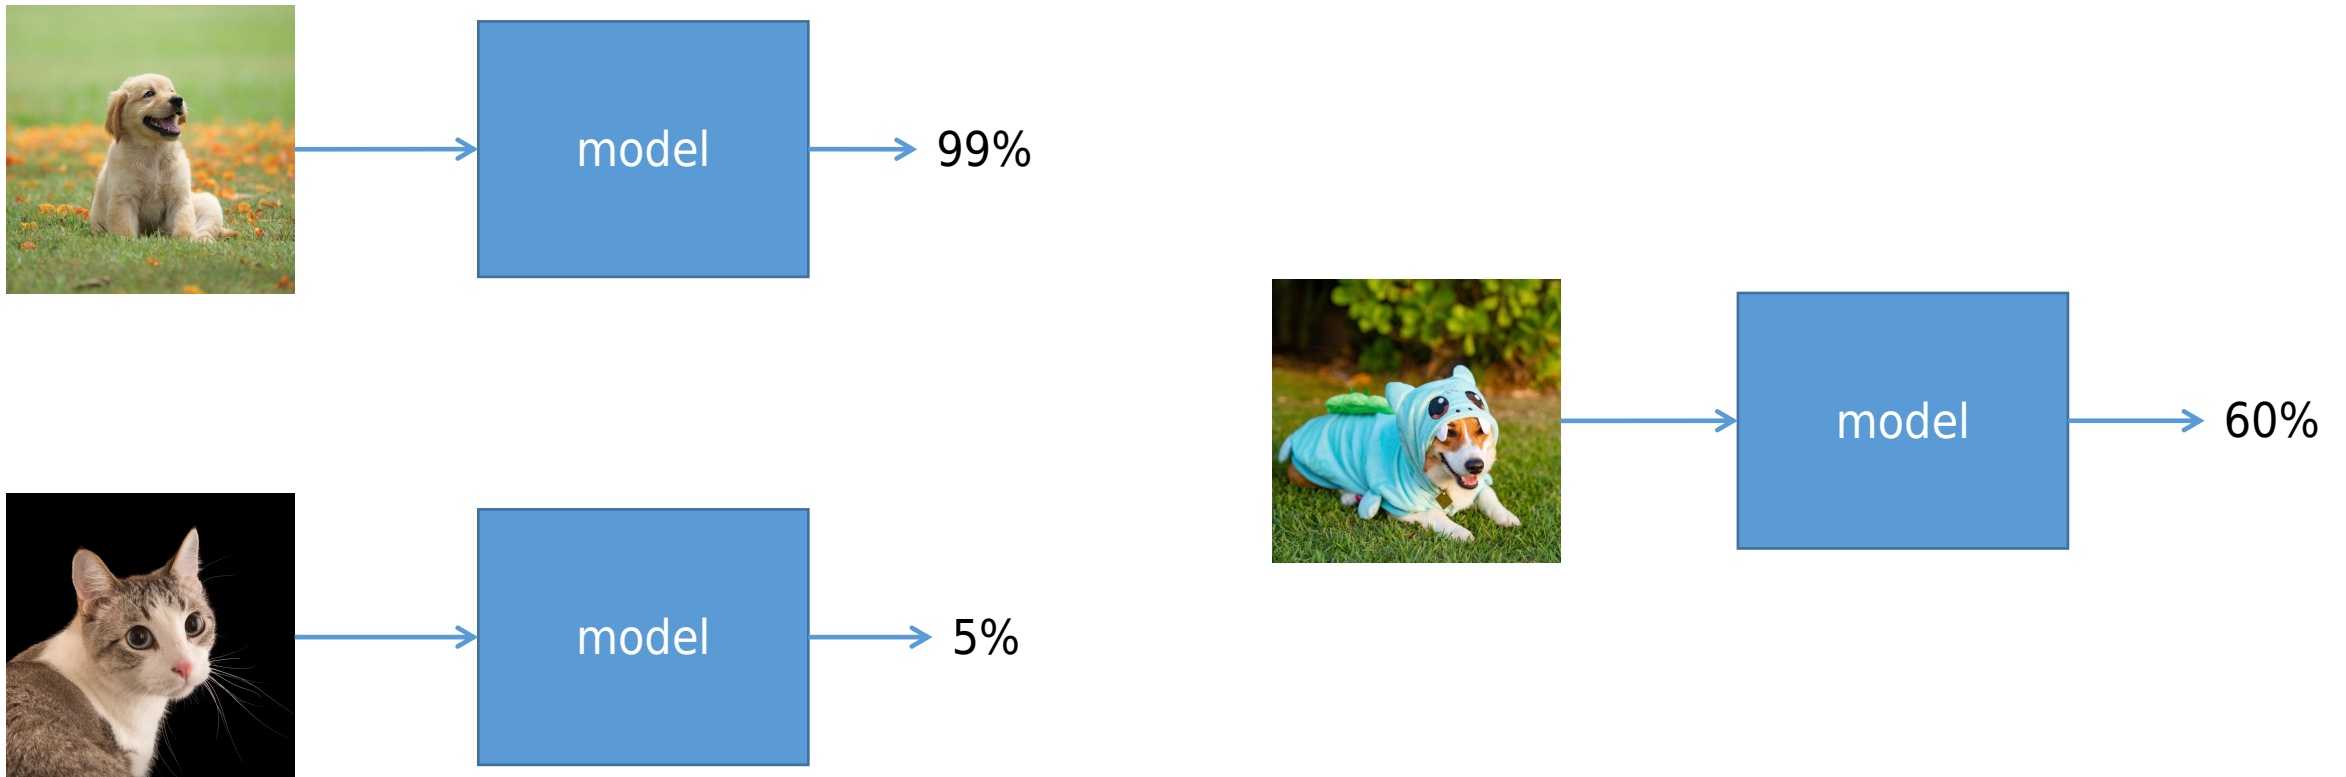
# Binary Classifier

A **binary classifier** is a **classifier** which can only **distinguish** between **two classes**.



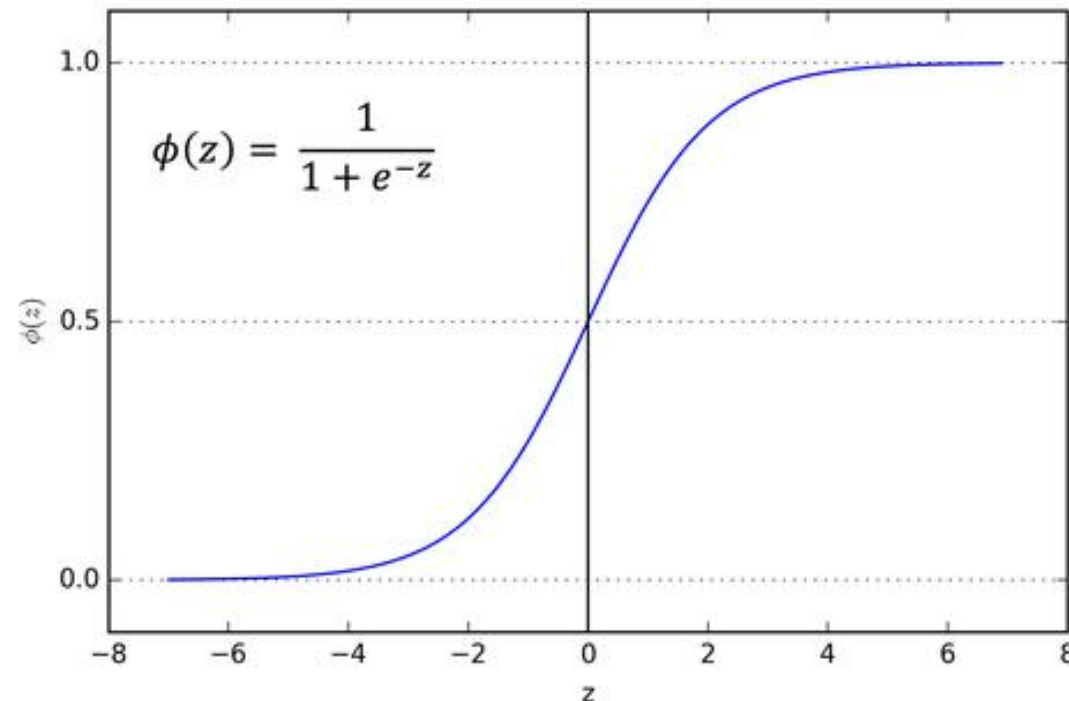
# Binary Classifier: Logistic Regression

In the **logistic regression** scenario, we aim to **train** our **network** to **output** the **probability** of the **input being** in the **target class**.



# Logistic Regression: Sigmoid Function

The **sigmoid function**, also **known** as the **logistic function**, is a function that **maps** any **real-valued number** to a value between **0** and **1**. It is commonly **used** in **logistic regression** to model the **probability** that a given **input belongs** to a particular **class**



## Binary Classifier: Logistic Regression loss

Given a **network output**  $o = model(x)$  and a **target class**  $t$  (which can be 0 or 1).

We can **transform**  $o$  **into a probability** using the **sigmoid** function  $\sigma(o)$ .

The **probability** of  $x$  of **being** of **class**  $t$  is then  $\sigma(o)$  and the **probability** of **not being** of **class**  $t$  is  $1 - \sigma(o)$ .

The Logistic Regression loss function is:

$$loss = -(t \cdot \log(\sigma(o)) + (1 - t) \cdot \log(1 - \sigma(o)))$$

Why this expression? why the minus sign? and why log?

## Binary Classifier: Logistic Regression loss

$$loss = -(t \cdot \log(\sigma(o)) + (1 - t) \cdot \log(1 - \sigma(o)))$$

If target class is 0 and the predicted probability  $\sigma(o)$  is 0.2

$$loss = -(0 \cdot \log(0.2) + 1 \cdot \log(1 - 0.2)) = -\log(0.8)$$

If target class is 1 and the predicted probability  $\sigma(o)$  is 0.8

$$loss = -(0 \cdot \log(0.8) + 1 \cdot \log(1 - 0.8)) = -\log(0.2)$$


# Binary Classifier: Logistic Regression loss

target 0 and predicted 0.8 =  $-\log(0.2)$

big error

target 0 and predicted 0.2 =  $-\log(0.8)$

small error



The error  
is exponentially  
larger w.r.t. the  
predicted  
probability  
error




## More than two classes

A **classifier** is **trained** to **classify** an **input** into **more** than two **classes**.

The **output** of the model **should** be a **categorical distribution**.

A **categorical distribution** is a probability distribution used to model the **likelihood** of **distinct classes**, **each class** has an associated **probability**.

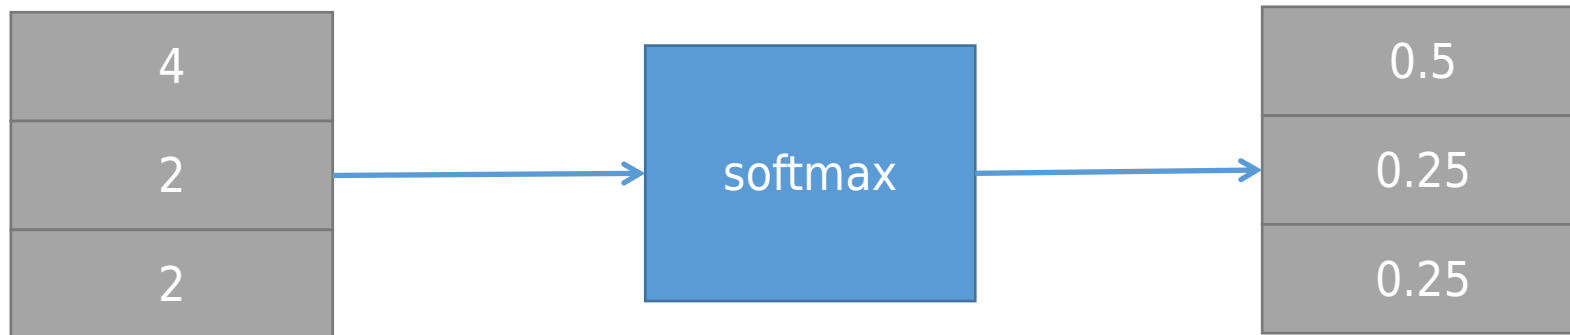
The **probabilities sum** to **1**, reflecting the **exclusive** nature of the **categories**.



# Classifier: the softmax function

The **softmax function** is a function that takes as **input** a **vector** of real **numbers** and transforms it **into** a **categorical probability** distribution.

$$\sigma(\mathbf{z})_i = \frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}}$$



# Classifier: the crossentropy loss

**Generalization** over **multiple classes** of the **logistic regression** loss:

$$loss = - \sum_{c=0}^n t_c \cdot \log(o_c)$$

where  $t$  is the one-hot encoded target class.

target = 2

0
0
1
0

target = 0

1
0
0
0

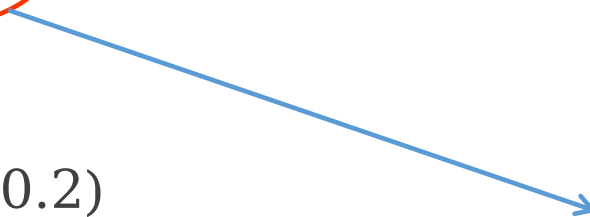
## Classifier: the crossentropy loss

$$loss = - \sum_{c=0}^n t_c \cdot \log(o_c)$$

Problem with 4 classes, target class is 1, output distribution is [0.1, 0.2, 0.2, 0.5]

$$t = [0 \ 1 \ 0 \ 0]$$

$$loss = -\log(0.2)$$



one-hot encoding  
of class 1

# Crossentropy loss in PyTorch

The **crossentropy** loss is implemented in **PyTorch** in the class **torch.nn.CrossEntropyLoss**.

The forward function of the **CrossEntropyLoss** module accepts a **tensor** in the **shape [batch, classes]** and a **tensor** of classes in the **shape [batch]**

```
import torch

loss_fn = torch.nn.CrossEntropyLoss()

logits = torch.randn(8, 5)
classes = torch.tensor([2, 1, 4, 2, 0, 3, 2, 0])

print(loss_fn(logits, classes))
```

```
import torch

loss_fn = torch.nn.NLLLoss()

logits = torch.randn(8, 5)
probs = torch.softmax(logits, dim=1)
log_probs = torch.log(probs)

classes = torch.tensor([2, 1, 4, 2, 0, 3, 2, 0])

print(loss_fn(log_probs, classes))
```